

```
In [1]: #dodda.prudhvi sai  
#assignment-04  
#simple linear regression
```

```
In [ ]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

```
In [2]: dataset =pd.read_csv('student_scores.csv')
```

```
In [3]: dataset
```

Out[3]:


	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25
10	7.7	85
11	5.9	62
12	4.5	41
13	3.3	42
14	1.1	17
15	8.9	95
16	2.5	30
17	1.9	24
18	6.1	67
19	7.4	69
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

```
In [4]: dataset.isnull().any()
```

```
Out[4]: Hours      False  
Scores      False  
dtype: bool
```

```
In [7]: plt.scatter(dataset['Scores'],dataset['Hours'])  
plt.xlabel('Scores')  
plt.ylabel('Hours')
```

```
Out[7]: Text(0, 0.5, 'Hours')
```



```
In [8]: #independent  
x=dataset.iloc[:,0:1]  
y=dataset.iloc[:,1:]
```

```
In [10]:  
x=dataset.iloc[:,0:1].values  
y=dataset.iloc[:,1:].values
```

```
In [11]: x.shape
```

```
Out[11]: (25, 1)
```

```
In [12]: y.shape
```

```
Out[12]: (25, 1)
```

```
In [13]: from sklearn.model_selection import train_test_split
```

```
In [14]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
In [15]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()
```

```
In [16]: lr.fit(x_train,y_train)
```

```
Out[16]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [17]: y_pred=lr.predict(x_test)
```

```
In [18]: y_pred
```

```
Out[18]: array([[16.88414476],
                [33.73226078],
                [75.357018   ],
                [26.79480124],
                [60.49103328]])
```

```
In [19]: from sklearn.metrics import r2_score
accuracy=r2_score(y_test,y_pred)
```

```
In [20]: accuracy
```

```
Out[20]: 0.9454906892105356
```


```
In [21]: y_pred_Scores=lr.predict([[15]])
```

```
In [22]: y_pred_Scores
```

```
Out[22]: array([[150.67800725]])
```

```
In [23]: plt.scatter(x_train,y_train,color='green')
plt.plot(x_test,lr.predict(x_test))
```

```
Out[23]: [<matplotlib.lines.Line2D at 0x1a263659f48>]
```



```
In [24]: lr.intercept_#intercept
```

```
Out[24]: array([2.01816004])
```

```
In [25]: lr.coef_#slope
```

```
Out[25]: array([[9.91065648]])
```

```
In [2]: #multi linear regression  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

```
In [3]: dataset=pd.read_csv('petrol_consumption.csv')
```

```
In [4]: dataset
```

20	7.00	4593	7834	0.663
21	8.00	4983	602	0.602
22	9.00	4897	2449	0.511
23	9.00	4258	4686	0.517
24	8.50	4574	2619	0.551
25	9.00	3721	4746	0.544
26	8.00	3448	5399	0.548
27	7.50	3846	9061	0.579
28	8.00	4188	5975	0.563
29	9.00	3601	4650	0.493
30	7.00	3640	6005	0.518

```
In [5]: type(dataset)
```

```
Out[5]: pandas.core.frame.DataFrame
```

```
In [7]: dataset.isnull().any()
```

```
Out[7]: Petrol_tax                False  
Average_income                False  
Paved_Highways                False  
Population_Driver_licence(%)  False  
Petrol_Consumption            False  
dtype: bool
```

In [8]: `dataset.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48 entries, 0 to 47
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Petrol_tax                            48 non-null     float64
1   Average_income                        48 non-null     int64
2   Paved_Highways                        48 non-null     int64
3   Population_Driver_licence(%)         48 non-null     float64
4   Petrol_Consumption                    48 non-null     int64
dtypes: float64(2), int64(3)
memory usage: 2.0 KB
```

In [10]: `dataset.head()`

Out[10]:

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)	Petrol_Consumption
0	9.0	3571	1976	0.525	541
1	9.0	4092	1250	0.572	524
2	9.0	3865	1586	0.580	561
3	7.5	4870	2351	0.529	414
4	8.0	4399	431	0.544	410

In [12]: `x=dataset.drop(['Petrol_Consumption'],axis=1).values`
`y=dataset['Petrol_Consumption'].values`

In [13]: `y`

Out[13]: `array([541, 524, 561, 414, 410, 457, 344, 467, 464, 498, 580, 471, 525, 508, 566, 635, 603, 714, 865, 640, 649, 540, 464, 547, 460, 566, 577, 631, 574, 534, 571, 554, 577, 628, 487, 644, 640, 704, 648, 968, 587, 699, 632, 591, 782, 510, 610, 524], dtype=int64)`

```
In [14]: x
```

```
Out[14]: array([[9.0000e+00, 3.5710e+03, 1.9760e+03, 5.2500e-01],
 [9.0000e+00, 4.0920e+03, 1.2500e+03, 5.7200e-01],
 [9.0000e+00, 3.8650e+03, 1.5860e+03, 5.8000e-01],
 [7.5000e+00, 4.8700e+03, 2.3510e+03, 5.2900e-01],
 [8.0000e+00, 4.3990e+03, 4.3100e+02, 5.4400e-01],
 [1.0000e+01, 5.3420e+03, 1.3330e+03, 5.7100e-01],
 [8.0000e+00, 5.3190e+03, 1.1868e+04, 4.5100e-01],
 [8.0000e+00, 5.1260e+03, 2.1380e+03, 5.5300e-01],
 [8.0000e+00, 4.4470e+03, 8.5770e+03, 5.2900e-01],
 [7.0000e+00, 4.5120e+03, 8.5070e+03, 5.5200e-01],
 [8.0000e+00, 4.3910e+03, 5.9390e+03, 5.3000e-01],
 [7.5000e+00, 5.1260e+03, 1.4186e+04, 5.2500e-01],
 [7.0000e+00, 4.8170e+03, 6.9300e+03, 5.7400e-01],
 [7.0000e+00, 4.2070e+03, 6.5800e+03, 5.4500e-01],
 [7.0000e+00, 4.3320e+03, 8.1590e+03, 6.0800e-01],
 [7.0000e+00, 4.3180e+03, 1.0340e+04, 5.8600e-01],
 [7.0000e+00, 4.2060e+03, 8.5080e+03, 5.7200e-01],
 [7.0000e+00, 3.7180e+03, 4.7250e+03, 5.4000e-01],
 [7.0000e+00, 4.7160e+03, 5.9150e+03, 7.2400e-01],
 [8.5000e+00, 4.3410e+03, 6.0100e+03, 6.7700e-01],
 [7.0000e+00, 4.5930e+03, 7.8340e+03, 6.6300e-01],
 [8.0000e+00, 4.9830e+03, 6.0200e+02, 6.0200e-01],
 [9.0000e+00, 4.8970e+03, 2.4490e+03, 5.1100e-01],
 [9.0000e+00, 4.2580e+03, 4.6860e+03, 5.1700e-01],
 [8.5000e+00, 4.5740e+03, 2.6190e+03, 5.5100e-01],
 [9.0000e+00, 3.7210e+03, 4.7460e+03, 5.4400e-01],
 [8.0000e+00, 3.4480e+03, 5.3990e+03, 5.4800e-01],
 [7.5000e+00, 3.8460e+03, 9.0610e+03, 5.7900e-01],
 [8.0000e+00, 4.1880e+03, 5.9750e+03, 5.6300e-01],
 [9.0000e+00, 3.6010e+03, 4.6500e+03, 4.9300e-01],
 [7.0000e+00, 3.6400e+03, 6.9050e+03, 5.1800e-01],
 [7.0000e+00, 3.3330e+03, 6.5940e+03, 5.1300e-01],
 [8.0000e+00, 3.0630e+03, 6.5240e+03, 5.7800e-01],
 [7.5000e+00, 3.3570e+03, 4.1210e+03, 5.4700e-01],
 [8.0000e+00, 3.5280e+03, 3.4950e+03, 4.8700e-01],
 [6.5800e+00, 3.8020e+03, 7.8340e+03, 6.2900e-01],
 [5.0000e+00, 4.0450e+03, 1.7782e+04, 5.6600e-01],
 [7.0000e+00, 3.8970e+03, 6.3850e+03, 5.8600e-01],
 [8.5000e+00, 3.6350e+03, 3.2740e+03, 6.6300e-01],
 [7.0000e+00, 4.3450e+03, 3.9050e+03, 6.7200e-01],
 [7.0000e+00, 4.4490e+03, 4.6390e+03, 6.2600e-01],
 [7.0000e+00, 3.6560e+03, 3.9850e+03, 5.6300e-01],
 [7.0000e+00, 4.3000e+03, 3.6350e+03, 6.0300e-01],
 [7.0000e+00, 3.7450e+03, 2.6110e+03, 5.0800e-01],
 [6.0000e+00, 5.2150e+03, 2.3020e+03, 6.7200e-01],
 [9.0000e+00, 4.4760e+03, 3.9420e+03, 5.7100e-01],
 [7.0000e+00, 4.2960e+03, 4.0830e+03, 6.2300e-01],
 [7.0000e+00, 5.0020e+03, 9.7940e+03, 5.9300e-01]])
```

```
In [15]: x.shape
```

```
Out[15]: (48, 4)
```

```
In [17]: y.shape
```

```
Out[17]: (48,)
```

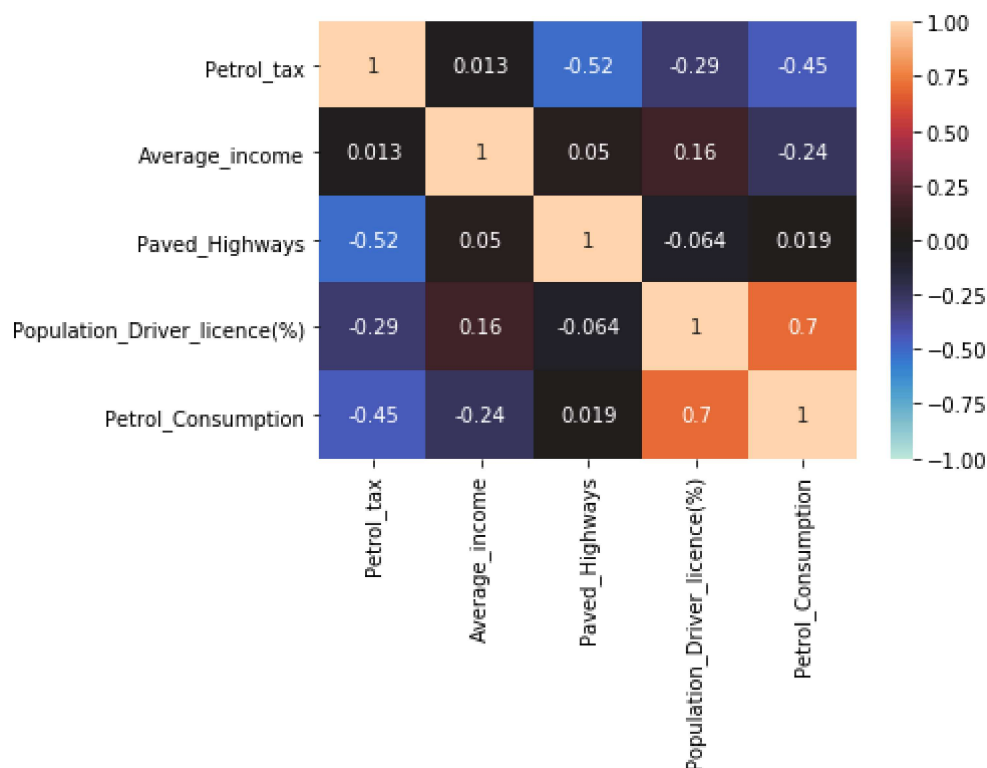
```
In [18]: dataset.corr()
```

```
Out[18]:
```

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence
Petrol_tax	1.000000	0.012665	-0.522130	-0.28
Average_income	0.012665	1.000000	0.050163	0.15
Paved_Highways	-0.522130	0.050163	1.000000	-0.06
Population_Driver_licence(%)	-0.288037	0.157070	-0.064129	1.00
Petrol_Consumption	-0.451280	-0.244862	0.019042	0.69

```
In [19]: import seaborn as sns
sns.heatmap(dataset.corr(),annot=True,vmin=-1,vmax=1,center=0)
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x15657e4b108>
```



```
In [35]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3, random_state=1)
```

```
In [36]: from sklearn.linear_model import LinearRegression
mlr=LinearRegression(normalize=True)
mlr.fit(x_train,y_train)
```

Out[36]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=True)

```
In [37]: y_pred=mlr.predict(x_test)
y_pred
```

Out[37]: array([620.58093461, 467.51687024, 739.21337689, 645.93715567,
591.85595105, 486.34203341, 583.80501833, 598.93055594,
708.77595273, 622.27676584, 565.00406531, 570.30820315,
571.26024921, 508.77150871, 702.78484107])

```
In [38]: mlr.predict([[9.0,3571,1976,0.525]])
```

Out[38]: array([534.03190111])

```
In [39]: from sklearn.metrics import r2_score
accuracy=r2_score(y_test,y_pred)
```

```
In [40]: accuracy
```

Out[40]: 0.6483616335976214

```
In [41]: mlr.intercept_
```


Out[41]: 304.6754038181706

```
In [42]: mlr.coef_
```

Out[42]: array([-2.85437070e+01, -7.36087171e-02, -1.65354087e-03, 1.43309331e+03])

```
In [45]: plt.scatter(y_test,y_pred)
```

Out[45]: <matplotlib.collections.PathCollection at 0x15658a58c08>



In []: